

EXPORTING MODELS AND INFERENCE IN KERAS

DR. MATTHEW SMITH

ADACS, SWINBURNE UNIVERSITY OF TECHNOLOGY

PREREQUISITES

- **WARNING:** This is not a stand-alone example.
- The testing and training data required to run these examples is not provided with this GIT repository.
- You will need the data contained within `ADACS_ML_A` and copy the training and test data into the Train and Test directories before use.
- If you need help, ask me or one of the other supervisors / instructors / gurus.

INFERENCE

- Inferencing is the process of:
 - Taking the learned “knowledge” of a previously trained machine, and
 - Application of this knowledge to new, previously unseen, data.
- As such, the process of inferencing might be considered as the goal of machine learning.
- Since there is no training (or heavy computation) the process of inferencing is usually quite fast.

INFERRNCING

- Up to now, the codes provided to you:
 - Create a Keras model,
 - Train it – perform computations on the training data set for computing the model weights (using `model.fit()`)
 - Test it – use the test data set to check the computed weights against known data sets (using `model.evaluate()`).

INFERCING

- To move forward, we need to:
 - Modify our existing codes to export our model in a form which can be loaded later on, and
 - Create a new code – which we shall call `infer.py` – which will perform inferencing given a single data set as an input.

EXPORTING YOUR KERAS MODEL

- Open your train.py file – the first few lines of code here should look familiar.
- The addition are the lines of code shown in the lower half.

```
# Fit the model using the training set
history = model.fit(X_train, Y_train, epochs=N_epochs, batch_size=32)

# Plot the history
plot_history(history)

# Final evaluation of the model using the Test Data
print("Evaluating Test Set")
scores = model.evaluate(X_test, Y_test, verbose=1)
print("Accuracy: %.2f%%" % (scores[1]*100))

# Export the model to file
model_json = model.to_json()
with open("model.json", "w") as json_file:
    json_file.write(model_json)
# Save the weights as well, in HDF5 format
model.save_weights("model.h5")
```

Existing

New

EXPORTING YOUR KERAS MODEL

- One method for exporting the model information to file is to employ the JSON format (pronounced JAY-SON)
- JSON: Javascript Object Notation
- This is simply a light-weight, text based format used for data exchange.
- The idea is that a JSON file is easily viewed by humans and interpreted by computer systems – it looks very similar to a C/C++ code.

```
{ "menu": {  
  "id": "file",  
  "value": "File",  
  "popup": {  
    "menuitem": [  
      { "value": "New", "onclick": "CreateNewDoc()" },  
      { "value": "Open", "onclick": "OpenDoc()" },  
      { "value": "Close", "onclick": "CloseDoc()" }  
    ]  
  }  
}  
}}
```

EXPORTING YOUR KERAS MODEL

- The object holding the json data is produced using the code:

```
model_json = model.to_json()
```

- This is then written to file for us to load later.

```
# Fit the model using the training set
history = model.fit(X_train, Y_train, epochs=N_epochs, batch_size=32)

# Plot the history
plot_history(history)

# Final evaluation of the model using the Test Data
print("Evaluating Test Set")
scores = model.evaluate(X_test, Y_test, verbose=1)
print("Accuracy: %.2f%%" % (scores[1]*100))

# Export the model to file
model_json = model.to_json()
with open("model.json", "w") as json_file:
    json_file.write(model_json)
# Save the weights as well, in HDF5 format
model.save_weights("model.h5")
```


EXPORTING YOUR KERAS MODEL

- This is a sample of the JSON file produced by Keras using the `to_json()` function:
(at the terminal type: `cat model.json <enter>`)

```
{
  "class_name": "Sequential",
  "keras_version": "2.1.4",
  "config": [
    {
      "class_name": "Dense",
      "config": {
        "kernel_initializer": {
          "class_name": "VarianceScaling",
          "config": {
            "distribution": "uniform",
            "scale": 1.0,
            "seed": null,
            "mode": "fan_avg"
          }
        },
        "name": "dense_1",
        "kernel_constraint": null,
        "bias_regularizer": null,
        "bias_constraint": null,
        "dtype": "float32",
        "activation": "relu",
        "trainable": true,
        "kernel_regularizer": null,
        "bias_initializer": {
          "class_name": "Zeros",
          "config": {}
        },
        "units": 16,
        "batch_input_shape": [null, 128],
        "use_bias": true,
        "activity_regularizer": null
      },
      "class_name": "Dense",
      "config": {
        "kernel_initializer": {
          "class_name": "VarianceScaling",
          "config": {
            "distribution": "uniform",
            "scale": 1.0,
            "seed": null,
            "mode": "fan_avg"
          }
        },
        "name": "dense_2",
        "kernel_constraint": null,
        "bias_regularizer": null,
        "bias_constraint": null,
        "activation": "softmax",
        "trainable": true,
        "kernel_regularizer": null,
        "bias_initializer": {
          "class_name": "Zeros",
          "config": {}
        },
        "units": 8,
        "use_bias": true,
        "activity_regularizer": null
      },
      "class_name": "Dense",
      "config": {
        "kernel_initializer": {
          "class_name": "VarianceScaling",
          "config": {
            "distribution": "uniform",
            "scale": 1.0,
            "seed": null,
            "mode": "fan_avg"
          }
        },
        "name": "dense_3",
        "kernel_constraint": null,
        "bias_regularizer": null,
        "bias_constraint": null,
        "activation": "softmax",
        "trainable": true,
        "kernel_regularizer": null,
        "bias_initializer": {
          "class_name": "Zeros",
          "config": {}
        },
        "units": 8,
        "use_bias": true,
        "activity_regularizer": null
      }
    ]
  }
}
```

EXPORTING YOUR KERAS MODEL

- In addition to this, we need to export the weights computed by the model.

`model.save_weights(INPUT)`

Where INPUT here is the name of the HDF5 file we wish to save to. This file will also be used later in infer.py.

```
# Fit the model using the training set
history = model.fit(X_train, Y_train, epochs=N_epochs, batch_size=32)

# Plot the history
plot_history(history)

# Final evaluation of the model using the Test Data
print("Evaluating Test Set")
scores = model.evaluate(X_test, Y_test, verbose=1)
print("Accuracy: %.2f%%" % (scores[1]*100))

# Export the model to file
model_json = model.to_json()
with open("model.json", "w") as json_file:
    json_file.write(model_json)

# Save the weights as well, in HDF5 format
model.save_weights("model.h5")
```

EXPORTING YOUR KERAS MODEL

- HDF5 File format = 5th generation Hierarchical Data Format (HDF), which is designed to store large amounts of data.
- Originally developed at the US's NCSA (National Center for Supercomputing Applications), home of the Blue Waters supercomputer (originally IBM, later Cray)
- We can load these files in Python (by importing the h5py module), but we won't need to do that directly.

INFER.PY – IMPORTING MODELS

IMPORTING YOUR KERAS MODEL

- Importing a previously generated Keras model is almost a simple reversal of the export steps – we need to:
 - Open the model.json file for reading and load the data,
 - Create a model from the information contained within the loaded data, and
 - Compile the model, in the same way we compiled the model during training.

IMPORTING YOUR KERAS MODEL

- Here are the key elements of this process.

```
# Load the JSON file
json_file = open('model.json','r')
loaded_model_json = json_file.read()
json_file.close()

# Set up the neural layer configuration in the model
model = model_from_json(loaded_model_json)
# Load the weights into the model
model.load_weights("model.h5")
# Compile it
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
```

- In this case, the model.json and model.h5 files are in the same directory as this python script.

IMPORTING YOUR KERAS MODEL

- Another note: here we are required to compile the model.
- This gives us an opportunity to use different optimizers and loss computation methods than those used during training. I wouldn't.

```
# Load the JSON file
json_file = open('model.json','r')
loaded_model_json = json_file.read()
json_file.close()

# Set up the neural layer configuration in the model
model = model_from_json(loaded_model_json)
# Load the weights into the model
model.load_weights("model.h5")
# Compile it
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
```

THE FINAL PRODUCT: INFER.PY

- Let's have a look at the entire code.
- The first new addition is the `model_from_json` module which needs to be imported.
- We can see this code parses the input provided from the command prompt – if the user does not provide an ID, we use `ID = 2` as a default value.

```
# infer.py
# Written by Dr. Matthew Smith, Swinburne University of Technol
# Load a precomputed keras model and its weights for a single i
# USAGE: python infer.py ID where ID is the ID of the training
# we wish to load.

# Import modules
import numpy as np
from keras.models import Sequential
from keras.models import model_from_json
from utilities import *
import sys

# Parse the input
no_arg = len(sys.argv)
if (no_arg == 2):
    ID = int(sys.argv[1])
else:
    print("Usage: python view.py <Data_ID>")
    print("where Data_ID is a number.")
    print("Example: python infer.py 2")
    print("  -- will load X_2.dat and infer its class")
    ID = 2

print("Loading file = " + str(ID))
```

THE FINAL PRODUCT: INFER.PY

- In this case, we are going to be lazy and load one of the test data sets for inferencing.
- Normally you would have the data you wished inferenced provided in another way – but since we are short on time, let's use the tools we have available.
- We use the `read_test_data()` function (from `utilities.py`) to load a single set of data in – if you wish to modify this later to load multiple sets, the tools are there for you.

```
# We still need to know how long the time series is
N_sequence = 128      # Length of each piece of data

# Create variables for use while inferencing.
# Keeping it in array form; you might want to inference
# multiple data sets later.
X_infer = np.empty([1,N_sequence])
Y_infer = np.empty(1)

# We can take this data from anywhere - let's load one of the training sets
X_infer[0,], Y_infer[0] = read_test_data(ID, N_sequence)
```


THE FINAL PRODUCT: INFER.PY

- After this, we are free to load our model (we could have done this first, but no matter).
- We use two functions to perform our inference:

```
# Load the JSON file
json_file = open('model.json','r')
loaded_model_json = json_file.read()
json_file.close()

# Set up the neural layer configuration in the model
model = model_from_json(loaded_model_json)
# Load the weights into the model
model.load_weights("model.h5")
# Compile it
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])

# Now try classifying the single data file we loaded
Class_infer = model.predict_classes(X_infer)

# Compute the class predictions - shouldn't be used as certainties.
Class_prob = model.predict(X_infer)

print("The predicted class is %d" % Class_infer[0])
print("Class Predictions: Class 0 = %f, Class 1 = %f" % ((1.0-Class_prob[0]), Class_prob[0]))
print("The actual loaded class is %d" % Y_infer[0])
```

THE FINAL PRODUCT: INFER.PY

- The purpose of our ML engine was to predict classes – i.e. perform a classification task.
- Hence we use the `predict_classes()` function – in this case, the function will return either a `[0]` or `[1]` – if we load more data sets, it will be an array of 0's or 1's.

```
# Now try classifying the single data file we loaded
Class_infer = model.predict_classes(X_infer)
```

- There are additional inputs for the `predict_classes` function: feel free to browse the Keras documentation for these.

THE FINAL PRODUCT: INFER.PY

- The previous function returned the predicted classes, which is very convenient for us.
- If we wish the raw output of the NN to be provided, we use the predict() function:

```
# Compute the class predictions – shouldn't be used as certainties.  
Class_prob = model.predict(X_infer)
```

- This function replaces the predict_proba() function from earlier versions of Keras, though predict_proba should still work if used here (and provide the same result)

THE FINAL PRODUCT: INFER.PY

predict

```
predict(x, batch_size=None, verbose=0, steps=None)
```

Generates output predictions for the input samples.

Computation is done in batches.

Arguments

- **x**: The input data, as a Numpy array (or list of Numpy arrays if the model has multiple inputs).
- **batch_size**: Integer. If unspecified, it will default to 32.
- **verbose**: Verbosity mode, 0 or 1.
- **steps**: Total number of steps (batches of samples) before declaring the prediction round finished. Ignored with the default value of `None`.

THE FINAL PRODUCT: INFER.PY

- Since we have a binary classification problem with a single output in our NN, we will have a single value returned.
- It's not really a probability – but just between you and me, let's pretend it is.

```
# Load the JSON file
json_file = open('model.json','r')
loaded_model_json = json_file.read()
json_file.close()

# Set up the neural layer configuration in the model
model = model_from_json(loaded_model_json)
# Load the weights into the model
model.load_weights("model.h5")
# Compile it
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])

# Now try classifying the single data file we loaded
Class_infer = model.predict_classes(X_infer)

# Compute the class predictions - shouldn't be used as certainties.
Class_prob = model.predict(X_infer)

print("The predicted class is %d" % Class_infer[0])
print("Class Predictions: Class 0 = %f, Class 1 = %f" % ((1.0-Class_prob[0]), Class_prob[0]))
print("The actual loaded class is %d" % Y_infer[0])
```

THE FINAL PRODUCT: INFER.PY

- Since we are using the functions contained within utilities.py, and loading the test data for inferencing, we should place infer.py in the same directories as train.py.
- When we call this script (python infer.py 24), this is what we get:

```
Loading file = 24
Loading file ./Test/X_24.dat
The predicted class is 0
Class Predictions: Class 0 = 0.931161, Class 1 = 0.068839
The actual loaded class is 0
```

- You can see that we are quite sure that the class is not class 1 (accurate), and we have correctly picked the class.

ACTIVITY

- Make sure the codes provided by GIT are working correctly. Don't forget to reload the modules required by ozstar if you haven't already (. script.sh)
- Make a copy of the previous directory (ADACS_ML_A) (cp -r ADACS_ML_A FOLDER_NAME) and attempt:
 - Modify train.py yourself to export the JSON information and weights data,
 - Write your own inference code which is called by:
python infer.py FILENAME

where FILENAME is the name of the file containing the data we wish to classify.